

Code for the model

compiler: Google Colab

Technologies used: CNN and GRU, Deep learning to test and train the dataset

Dataset used: <https://drive.google.com/uc?id=1YlvpDLix3S-U8fd-gqRwPcWXAXm8JwjL>

```
!pip list
```

```
!pip install opencv-python matplotlib imageio gdown tensorflow
```

```
import os
```

```
import cv2
```

```
import tensorflow as tf
```

```
import numpy as np
```

```
from typing import List
```

```
from matplotlib import pyplot as plt
```

```
import imageio
```

```
tf.config.list_physical_devices('GPU')
```

```
physical_devices = tf.config.list_physical_devices('GPU')
```

```
try:
```

```
    tf.config.experimental.set_memory_growth(physical_devices[0], True)
```

```
except:
```

```
    pass
```

```
import gdown
```

```
url = 'https://drive.google.com/uc?id=1YlvpDLix3S-U8fd-gqRwPcWXAXm8JwjL'
```

```
output = 'data.zip'
```

```
gdown.download(url, output, quiet=False)
```

```
gdown.extractall('data.zip')
```

```
def load_video(path:str)->List[float]:
```

```
    cap=cv2.VideoCapture(path)
```

```
    frames=[]
```

```
    for _ in range(int(cap.get(cv2.CAP_PROP_FRAME_COUNT))):
```

```
        ret,frame=cap.read()
```

```
        frame=tf.image.rgb_to_grayscale(frame)
```

```

    # to extract lip
    # dlip used to extract part using that we get portion measurements
    frames.append(frame[190:236,80:220,:])
cap.release()
mean=tf.math.reduce_mean(frames)
std=tf.math.reduce_std(tf.cast(frames,tf.float32))
return tf.cast((frames-mean),tf.float32)/std
# each character we expect in video
vocab= [x for x in "abcdefghijklmnopqrstuvwxyz'?!123456789 "]
print(vocab)

char_to_num=tf.keras.layers.StringLookup(vocabulary=vocab,oov_token="")
# takes number convert to number
num_to_char=tf.keras.layers.StringLookup(
    vocabulary=char_to_num.get_vocabulary(),oov_token="",invert=True
)
print(
    f"The vocabulary is:{char_to_num.get_vocabulary()} "
    f"(size={char_to_num.vocabulary_size()})"
)
# this convert to number
print(char_to_num(['a','b','c']))

# to load alignments
def load_alignments(path:str)->List[str]:
    with open(path,'r') as f:
        lines=f.readlines()
    tokens =[]
    for line in lines:
        line=line.split()
        if line[2]!='sil':
            tokens=[*tokens,' ',line[2]]
    return char_to_num(tf.reshape(tf.strings.unicode_split(tokens,input_encoding='UTF-
8'),(-1)))[1:]

def load_data(path:str):

```

```

path=bytes.decode(path.numpy())
file_name=path.split('/')[-1].split('.')[0]
video_path=os.path.join('data','s1',f'{file_name}.mpg')
alignment_path=os.path.join('data','alignments','s1',f'{file_name}.align')
frames=load_video(video_path)
alignments=load_alignments(alignment_path)
return frames,alignments

```

```
test_path='./data/s1/bbal6n.mpg'
```

```

# test_path='./data\\s1\\bbal6n.mpg'
print(tf.convert_to_tensor(test_path).numpy().decode('utf-8').split('\\')[-1].split('.')[0])
print(tf.convert_to_tensor(test_path).numpy().decode('utf-8').split('\\'))
frames,alignments=load_data(tf.convert_to_tensor(test_path))
print(frames)
print(plt.imshow(frames[0]))
plt.imshow(frames[0])

```

```

print(alignments)
print(num_to_char(alignments.numpy()))

```

```

print([bytes.decode(x) for x in num_to_char(alignments.numpy()).numpy()])
print(tf.strings.reduce_join([bytes.decode(x) for x in
num_to_char(alignments.numpy()).numpy()])))

```

```

import tensorflow as tf
from typing import Tuple # Import the Tuple type

```

```

# Define your mappable_function to load data
def mappable_function(path:str) -> Tuple[tf.Tensor, tf.Tensor]:
    result = tf.py_function(load_data, [path], (tf.float32, tf.int64))
    return result

```

```

from matplotlib import pyplot as plt
data=tf.data.Dataset.list_files('./data/s1/*.mpg')
data=data.shuffle(500)
# paths of each file
data=data.map(mappable_function)
# print(frames)
# alignments of each file under dataset
print(alignments)

data=data.padded_batch(2,padded_shapes=([75,None,None,None],[40]))
data=data.prefetch(tf.data.AUTOTUNE)
frames,alignments = data.as_numpy_iterator().next()
# print(alignments)
test=data.as_numpy_iterator()
val=test.next();val[0]

sample = data.as_numpy_iterator()
val = sample.next()
fv = (val[0][0].astype(np.uint8) * 255).squeeze()
# converts into gif what he is saying
imageio.mimsave('./animation.gif', fv, duration=10)

from keras.models import *
from keras.layers import *
from keras.optimizers import Adam
from keras.callbacks import ModelCheckpoint, LearningRateScheduler

print(data.as_numpy_iterator().next()[0][0].shape)

```

```

model = Sequential()
model.add(Conv3D(128, 3, input_shape=(75,46,140,1), padding='same'))
model.add(Activation('relu'))
model.add(MaxPool3D((1,2,2)))

model.add(Conv3D(256, 3, padding='same'))
model.add(Activation('relu'))
model.add(MaxPool3D((1,2,2)))

model.add(Conv3D(75, 3, padding='same'))
model.add(Activation('relu'))
model.add(MaxPool3D((1,2,2)))

model.add(TimeDistributed(Flatten()))

model.add(Bidirectional(LSTM(128, kernel_initializer='Orthogonal',
return_sequences=True)))
model.add(Dropout(.5))

model.add(Bidirectional(LSTM(128, kernel_initializer='Orthogonal',
return_sequences=True)))
model.add(Dropout(.5))

model.add(Dense(char_to_num.vocabulary_size()+1, kernel_initializer='he_normal',
activation='softmax'))

model.summary()
yhat = model.predict(val[0])
yhat[0].shape
print(yhat)
print(tf.argmax(yhat[0],axis=1))

print(tf.strings.reduce_join([num_to_char(tf.argmax(x)) for x in yhat[0]]))
print(tf.strings.reduce_join([num_to_char(x) for x in tf.argmax(yhat[0],axis=1)]))

```

```
print( model.input_shape)
print(model.output_shape)
```

```
# setup training options and train
```

```
def scheduler(epoch,lr):
```

```
    if epoch<30:
```

```
        return lr
```

```
    else:
```

```
        return lr*tf.math.exp(-0.1)
```

```
def CTCLoss(y_true, y_pred):
```

```
    batch_len = tf.cast(tf.shape(y_true)[0], dtype="int64")
```

```
    input_length = tf.cast(tf.shape(y_pred)[1], dtype="int64")
```

```
    label_length = tf.cast(tf.shape(y_true)[1], dtype="int64")
```

```
    input_length = input_length * tf.ones(shape=(batch_len, 1), dtype="int64")
```

```
    label_length = label_length * tf.ones(shape=(batch_len, 1), dtype="int64")
```

```
    loss = tf.keras.backend.ctc_batch_cost(y_true, y_pred, input_length, label_length)
```

```
    return loss
```

```
class ProduceExample(tf.keras.callbacks.Callback):
```

```
    def __init__(self, dataset) -> None:
```

```
        self.dataset = dataset.as_numpy_iterator()
```

```
    def on_epoch_end(self, epoch, logs=None) -> None:
```

```
        data = self.dataset.next()
```

```
        yhat = self.model.predict(data[0])
```

```
        decoded = tf.keras.backend.ctc_decode(yhat, [75, 75], greedy=False)[0][0].numpy()
```

```
        for x in range(len(yhat)):
```

```
            print('Original:',
```

```
tf.strings.reduce_join(num_to_char(data[1][x])).numpy().decode('utf-8'))
```

```
            print('Prediction:',
```

```
tf.strings.reduce_join(num_to_char(decoded[x])).numpy().decode('utf-8'))
```

```
            print('~' * 100)
```

```

model.compile(optimizer=Adam(learning_rate=0.0001), loss=CTCLoss)
checkpoint_callback = ModelCheckpoint(os.path.join('models','checkpoint'),
monitor='loss', save_weights_only=True)
schedule_callback = LearningRateScheduler(scheduler)
example_callback = ProduceExample(data)
model.fit(data,epochs=100, callbacks=[checkpoint_callback,
schedule_callback,example_callback])

# to make prediction
import gdown
url='https://drive.google.com/uc?id=1vWscXs4Vt0a_1IH1-ct2TCgXAZT-N3_Y'
output='checkpoints.zip'
gdown.download(url,output,quiet=False)
gdown.extractall('checkpoints.zip','models')
model.load_weights('models/checkpoint')
test_data=data.as_numpy_iterator()
test_data.next()
sample=test_data.next()
yhat=model.predict(sample[0])
print('~'*100, 'REAL TEXT')
[tf.strings.reduce_join([num_to_char(word) for word in sentence]) for sentence in
sample[1]]

decoded=tf.keras.backend.ctc_decode(yhat,input_length=[75,75],greedy=True)[0][0].nu
mpty()

print('~'*100,'PREDICTIONS')
[tf.strings.reduce_join([num_to_char(word) for word in sentence]) for sentence in
decoded]
# test on a video

sample=load_data(tf.convert_to_tensor('.\\data\\s1\\prac6n.mpg'))
yhat=model.predict(tf.expand_dim(sample[0],axis=0))
# print('~'*100,'REAL TEXT')
print('~'*100, 'REAL TEXT')
[tf.strings.reduce_join([num_to_char(word) for word in sentence]) for sentence in
[sample[1]]]

```

```
decoded=tf.keras.backend.ctc_decode(yhat,input_length=[75],greedy=True)[0][0].numpy  
()
```

```
print('~'*100,PREDICTIONS')
```

```
[tf.strings.reduce_join([num_to_char(word) for word in sentence]) for sentence in  
decoded]
```